Tools For Data Science: Optimizing Team USA's Gymnastics Roster for the Paris 2024 Olympics

Rami Pellumbi[*]

December 20, 2023

[*] M.S., Statistics & Data Science

1 Introduction

In the high-stakes realm of Olympic gymnastics, selecting the optimal team composition is crucial for national success. This study presents a sophisticated system designed to assist Team USA in formulating their Men's and Women's Artistic Gymnastics teams for the Paris 2024 Olympics. Addressing the complex dynamics of Olympic competition, our system employs a blend of statistical modeling and advanced application design. At its core, the system utilizes linear regression models tailored to each gender-apparatus pair, predicting athletes' execution scores (`e_score`) based on their difficulty scores (`d_score`) and `name`. This modeling approach leverages data from the USCAS 2024 Gymnastics Data Challenge, ensuring robust and relevant insights. Notably, the `d_score` is used as a constant feature for an athlete on a given apparatus, allowing for better model performance on the prediction `e_score`. The performance of these models is acceptable, explaining reasonable amounts of the variance for gender-apparatus pairs. Model efficacy is achieved through meticulous data processing and model tuning. To translate these insights into actionable strategies, the system conducts parallel Monte Carlo simulations of the Olympic events. These simulations project a range of outcomes, offering medal counts for team medals, individual all-arounds, and apparatus-specific competitions.

2 Introduction

The difficulty in selecting the optimal team composition for the Olympics is a common theme across many sports. In gymnastics, this challenge is accentuated by the complex dynamics of Olympic competition. Each team present sends 5 athletes. For a team of 5 athletes, only 4 athletes can compete on each apparatus. This constraint yields $\binom{5}{4}^6$ and $\binom{5}{4}^4$ possible apparatus assignments for men and women, respectively. Combined with the fact that *each country* may send a different group of 5 and assign to apparatuses in an unpredictable manner, the number of possible team compositions is astronomical. Simulating all possible team compositions is computationally infeasible, requiring a more sophisticated approach. Our system addresses this challenge by empowering Team USA with a tool they can use to answer key questions, giving coaches and decision makers a data-driven approach to team selection.

Since there are so many possible combinations, the idea of what a "good" team composition is can be subjective. For example, a team that wins the most apparatus medals may not be the best team overall. Depending on the goal, the choice of team composition may vary. For this reason, our system produces all outcomes for a selected team, allowing the user to assess and weight the outcomes based on their preferences.

We start by building a linear regression model that predicts athletes' execution scores (`e_score`) based on their difficulty scores (`d_score`) and `name`. We justify the use of `d_score` as a feature to the model in Section 3. From there, we turn to Sections 4 and 5, where we discuss the model, assumptions, parameters, performance, and application features. In Section 5, we discuss the system architecture and how the model is integrated with the client application. Finally, we conclude with a discussion of the system's intended use case and next steps in Section 6.

3 Data Exploration

We assess the viability of the predictors of our model and explore the relationship between the predictors and the response variable. We then discuss the problems in the dataset and how we addressed them.

3.1 Model Viability

To assess the viability of using `d_score` as a feature in our model, we consider the four assumptions of the linear regression model $Y + X\beta + \epsilon$:

1. Linearity: The relationship between X and the mean of Y is linear.

2. Independence: Observations are independent of each other.

3. Homoscedasticity: The variance of residuals is the same for any value of X.
4. Normality: For a fixed value of X, Y is normally distributed.

3.1.1 Linearity

We assess the linearity assumption by plotting `e_score` against `d_score`. Figure 1. We observe that with the exception of a few outliers, the relationship between the mean of `e_score` and `d_score` is approximately linear.

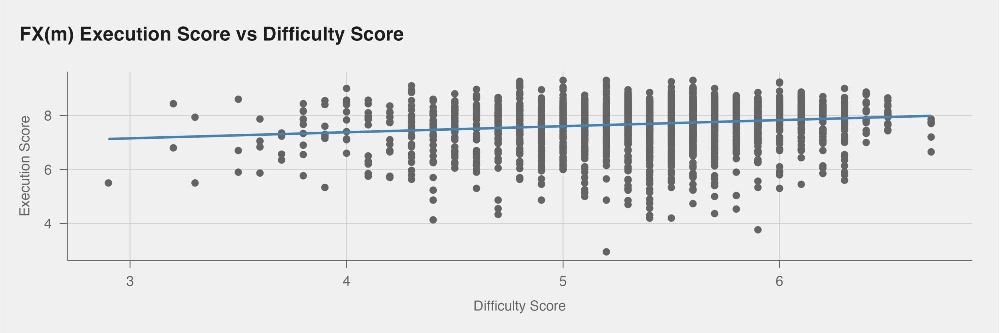

Figure 1: Execution Score vs Difficulty Score for FX (m)

3.1.2 Independence

The independence assumption is the easiest to verify. Since the data is collected from different athletes, we can assume that the observations are independent of each other.

3.1.3 Homoscedasticity

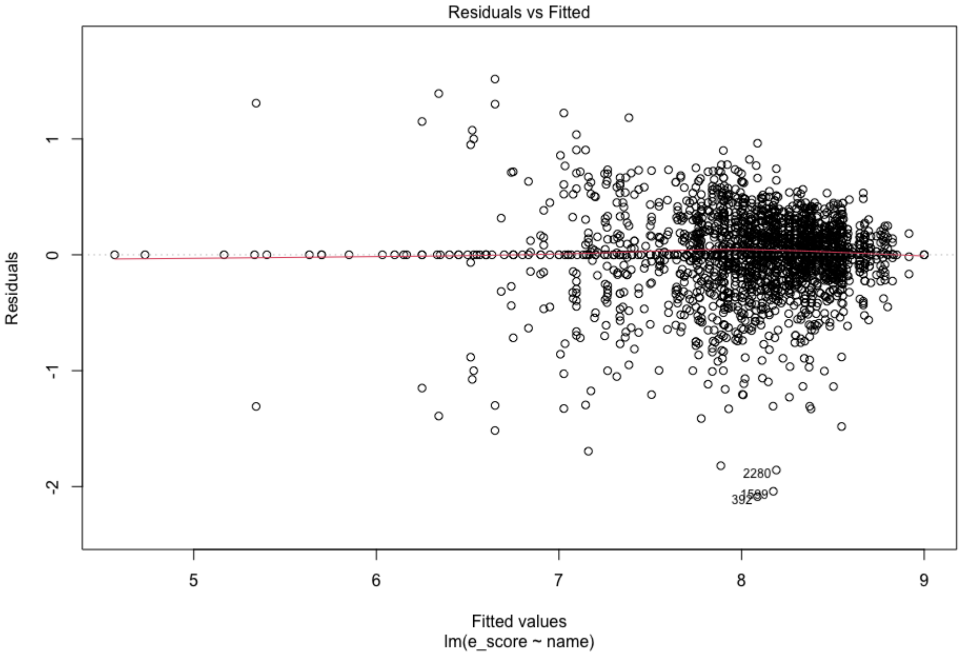

Figure 2: Residuals vs. Fitted for SR (m)

The residuals vs. fitted plot in Figure 2 shows that the residuals are roughly spread equally around the horizontal line without forming any discernible pattern. We conclude that the homoscedasticity assumption is satisfied.

3.1.4 Normality

We plot the distribution of `e_score` at different levels of `d_score`. We see that the distribution of `e_score` is approximately normal, albeit left-skewed, for the majority of `d_score` values. For outlier values, the distribution of `e_score` does not conform to the normality assumption. However, since these values are outliers and infrequent, we proceed with the assumption that the distribution of `e_score` is normal for a fixed value of `d_score`. Figures 3 and 4 show the distribution of `e_score` for VT (w) and PH (m), respectively.

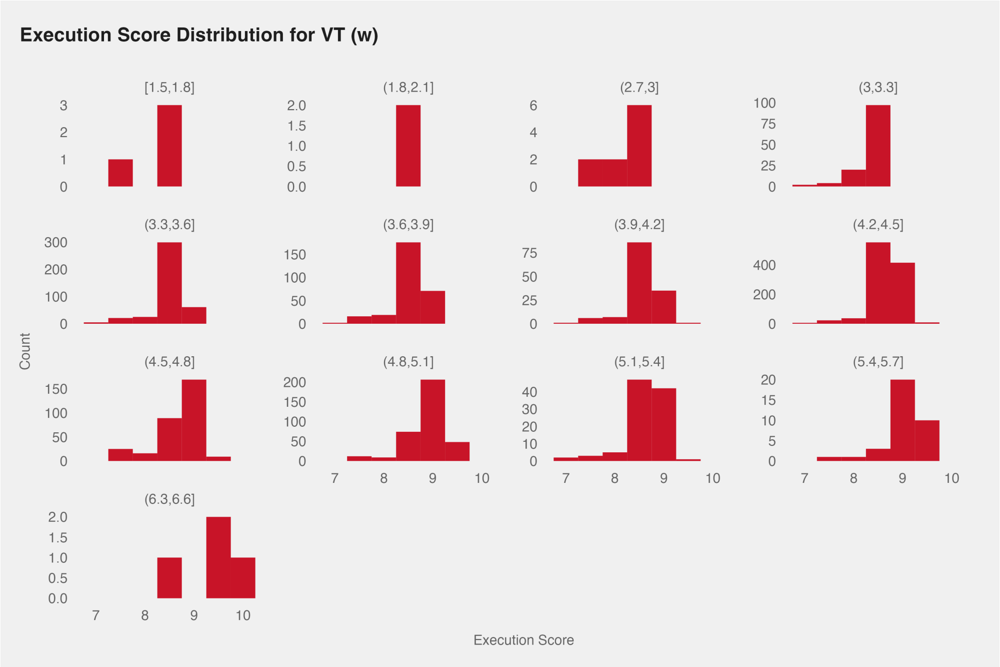

Figure 3: Distribution of `e_score` for VT (w)

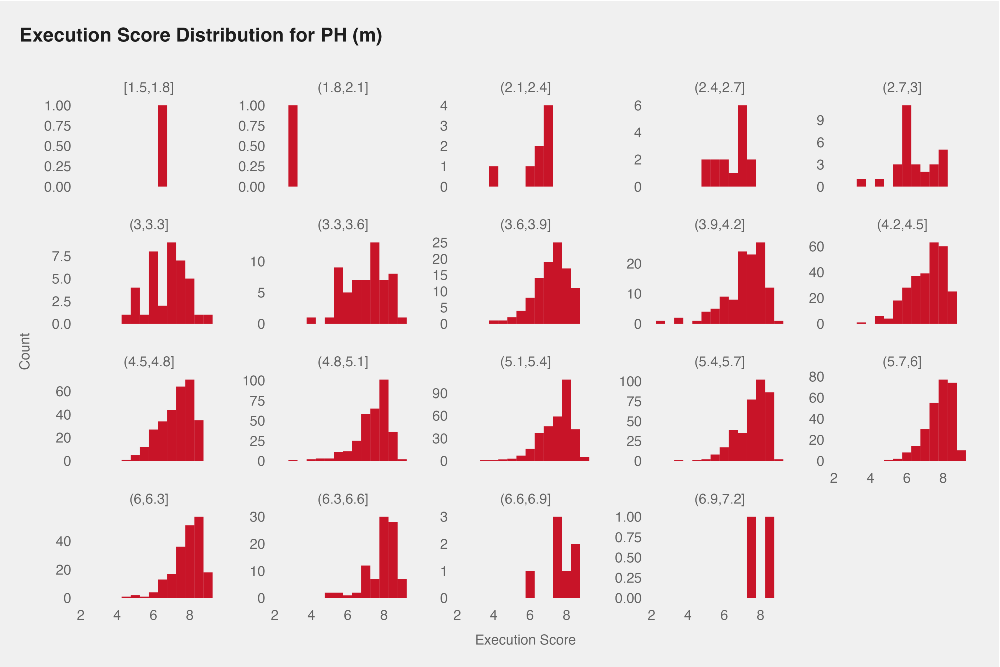

Figure 4: Distribution of `e_score` for PH (m)

3.2 Date Cleanup

The data was found to be malformed in several ways. First, there were many repeat names in the data with slight variations. That coupled with incorrect country codes and multiple countries associated to the same athlete forced us to clean the data. All duplicate athlete names were removed, and the country codes were corrected. Athletes with multiple countries were assigned the country with the most occurrences, except for athletes from Northern Ireland, which were forced to be part of NIR rather than GBR. Moreover, the data columns and values were made uniform, e.g., all lowercase or uppercase. Lastly, there is infrastructure in place to quickly swap out the randomly chosen alternates for the actual qualifiers once the Olympic team is announced. For alternates that have been announced, they are included in future simulations.

4 Modeling

4.1 Model Selection

The linear model $Y = X\beta + \epsilon$ was chosen for its simplicity and interpretability. The model is easy to understand and explain, which is important for a system that is intended to be used by coaches and decision-makers. Moreover, the model is easy to implement and computationally efficient for simulations. The exact model formula used was `e_score` $\sim$ `d_score` $+$ `name`. A different model is used for each apparatus-gender pair, giving us 10 models in total.

4.1.1 Modeling Parameters

The model parameter interpretations follow the standard linear regression model. For categorical features like name, the coefficient represents the expected change in `e_score` when the name is 'Athlete X' as opposed to 'Athlete Y' (the reference athlete left out by the model). For numerical features like `d_score`, the coefficient represents the expected change in `e_score` when `d_score` increases by one unit. Due to the complexity of model layouts, we do not include the model parameter values in this report.

4.1.2 Model Performance

The model performance as measured by R^2 is shown in figures 5 and 6. The models explain a reasonable amount of the variance and have acceptable performance.

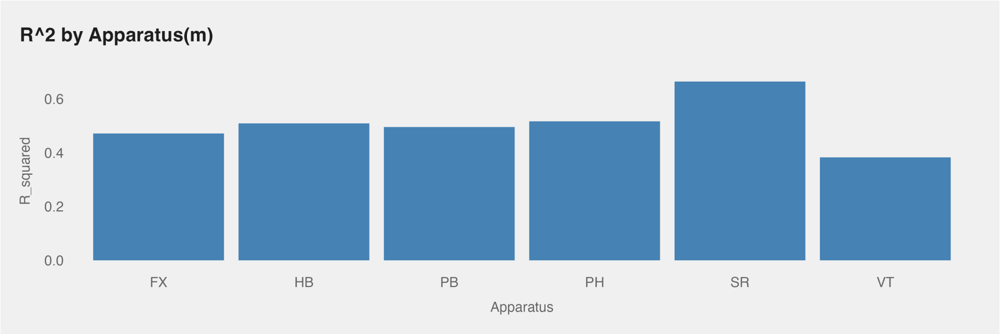

Figure 5: R^2 for FX (m)

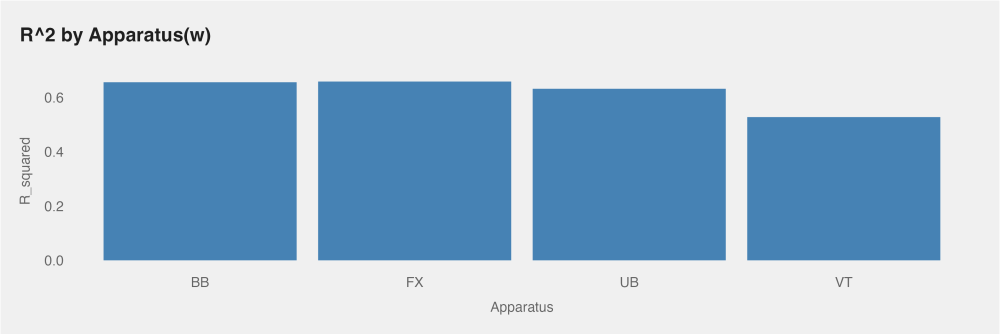

Figure 6: R^2 for FX (w)

The best model is chosen via the AIC criterion from the `MASS` package.

4.2 Simulations

We explore the simulation strategy. The only fixed competitors to the competition are the named competitors qualifying via criterias 4-7. For the non USA qualifying teams, we select the top 5 all around predicted scores for a country as the team of 5. For example, we run one round of our model for each country, and then sum up the predicted scores for each individual across all apparatuses. The top 5 for that country are selected to be the team of 5 for that country. The top 5 for team USA is omitted as this is controlled via the simulations.

There are three levels to simulations: team USA assignment, apparatus allocations, and number of simulations to run. For each team of 5 for team USA, we randomly assign groups of 4 to each apparatus for each country. For these fixed allocations, we run multiple simulations to obtain a medal count distribution. The medal count is determined according to the Olympic rules for progression in the competition. For example, the top 8 teams in the team all around competition advance to the team finals. The top 24 individuals in the individual all around competition advance to the individual all around finals, with a maximum of two per country, etc.

5 System Architecture

This application follows the three-tier architecture. That is, there is a client application that acts only as a view layer that does not contain any direct database calls. The client communicates with the server through a forms interface, and the server in turn communicates with a database system to access data. More details on the system architecture are given in the repository README files.

5.1 Server & Database

The server and database are both implemented in R. The server is a Plumber API that is run via RStudio's `plumber` package, while the database is a sequence of RDS files loaded by the API on startup. For more details on the server and R code, see the repository's server README file.

5.1.1 Server

The server is a Plumber API that is run via RStudio's `plumber` package. It is responsible for handling requests from the client and returning the appropriate response. In pariticular, the server is able to explore simulations run offline as well as run new simulations in real time.

5.1.2 Database

The 'database' is a sequence of RDS files loaded by the API on startup. The RDS files are structured as one to many relationships, where each file contains relevant simulation data. The database files are created by following the server's README file.

5.2 Client Application

The client application consists of two pages: the simulation runner and the simulation explorer. The simulation runner allows the user to run new simulations and view the results. The simulation explorer allows the user to explore simulations that have already been run. We developed a simple user interface that allows the user to select a team of 5, allocate them to apparatuses, and run simulations. The results are displayed in the form of bar charts, showing the medal count for each team or individual through the simulations.

5.2.1 Simulation Runner

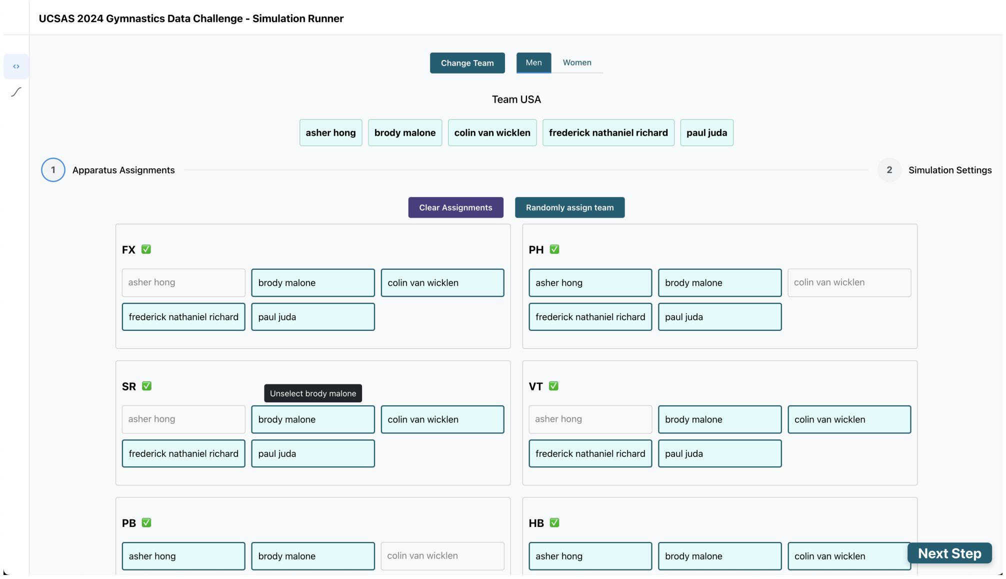

UCSAS 2024 Gymnastics Data Challenge - Simulation Runner

Team USA

asher hong | brody malone | colin van wicklen | frederick nathaniel richard | paul juda

FX ✅: asher hong, brody malone, colin van wicklen, frederick nathaniel richard, paul juda

PH ✅: asher hong, brody malone, colin van wicklen, frederick nathaniel richard, paul juda

SR ✅: asher hong, brody malone, colin van wicklen, frederick nathaniel richard, paul juda

VT ✅: asher hong, brody malone, colin van wicklen, frederick nathaniel richard, paul juda

PB ✅: asher hong, brody malone, colin van wicklen

HB ✅: asher hong, brody malone, colin van wicklen

Figure 7: Simulation Runner Apparatus Allocations

The simulation runner, shown in Figure 7, allows the user to select a gender and associated team of 5, from which they allocate to apparatuses. The user is able to do the assignments manually or pick one at random.

There are a variety of UI features that prevent the user from picking a bad team, e.g., an apparatus allocation is only possible for users that have data points for that apparatus. If no valid team is assignable, the user is notified. Once a valid team is selected, the user sees results as in Figure 8. A figure is shown for team all around, individual all around, and all apparatus events.

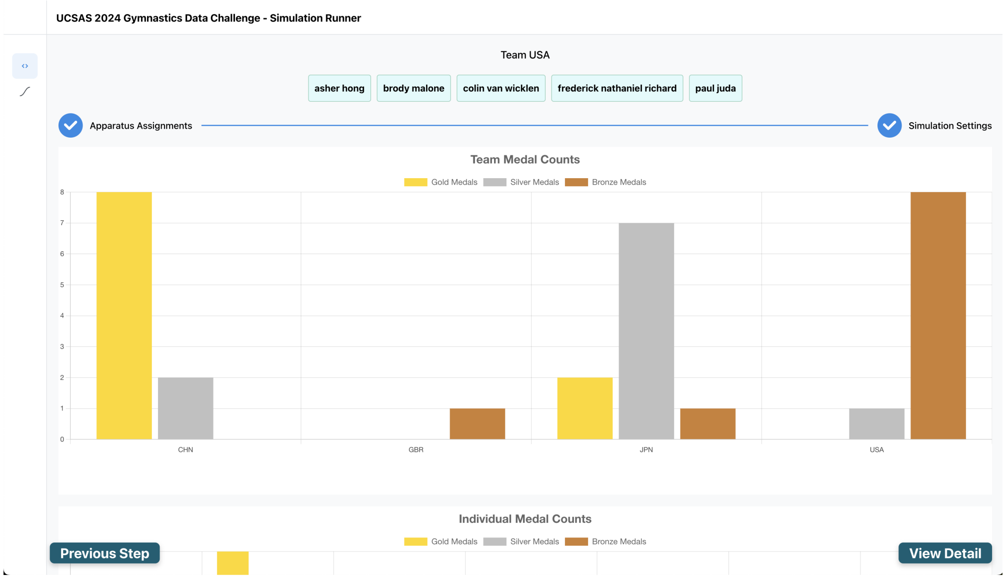
UCSAS 2024 Gymnastics Data Challenge - Simulation Runner

Team USA

asher hong, brody malone, colin van wicklen, frederick nathaniel richard, paul juda

Figure 8: Simulation Runner Results

Additionally, the user is able to view the apparatus and team allocations of all other countries in the simulations. This allows the user to compare their team to other teams and see how they stack up.

5.2.2 Simulation Explorer

The simulation explorer runs entirely similar to the simulation runner, except that the user is able to explore simulations that have already been run via the `run_simulations.R` script in the server. This process is detailed in the repository.

5.2.3 Implications

The system offers a data-driven approach to team selection, allowing coaches and decision-makers to make informed choices. The idea was not to provide recommendations, it was to provide a tool to non data scientists to let them answer their own questions. The user experience is designed to be simple and intuitive, allowing for easy exploration of different team compositions and outcomes. The choice of bar charts for the results was intentional, as it allows for easy comparison of different teams and individuals across a variety of simulation.

6 Conclusion

This report has presented a comprehensive system for optimizing Team USA's gymnastics roster for the Paris 2024 Olympics. The initiative leverages advanced statistical models and simulation techniques to aid in making informed decisions about team composition. In this conclusion, we reflect on the system's key features, its implications, and propose future enhancements.

The core of the system is a set of linear regression models, each tailored to a specific gender-apparatus pair. These models predict athletes' execution scores based on their difficulty scores and names. Our thorough data exploration and cleansing process ensure the integrity and reliability of these models. We have addressed potential issues such as data inconsistencies and athlete name duplications, resulting in a robust dataset.

The modeling approach, rooted in simplicity and interpretability, is a significant strength of the system. By focusing on linear regression, we facilitate ease of understanding and implementation, crucial for stakeholders such as coaches and decision-makers. Our models perform satisfactorily across different apparatuses and genders, as evidenced by reasonable R^2 values, implying a good fit for the data.

The simulation component of the system, which employs Monte Carlo methods, is particularly notable. It allows for the exploration of various team compositions and outcomes, reflecting the complexity and unpredictability of Olympic competitions. This aspect is crucial in a sport where strategic team composition can significantly impact the overall performance and medal prospects.

The system's primary utility lies in its ability to provide a data-driven approach to team selection, a task historically influenced by subjective judgment and experience. By quantifying performance predictions and outcomes, the system empowers coaches and decision-makers with actionable insights by enabling a strategic evaluation of different team compositions.

Moreover, the system's architecture and user interface are designed for accessibility and ease of use. The simulation runner and explorer facilitate interactive exploration of potential outcomes, a feature that is particularly valuable for strategic planning and scenario analysis.